

To achieve a reliable solution for an EMA (Enterprise Market Access) consumer application running on two Linux hosts with one active and one inactive session, ensuring high availability and failover, several coordination approaches can be used. Each approach must ensure that only one host is active at a time, with the inactive host taking over if the active host fails (e.g., due to a crash or connection loss detected via `StatusMsg`). The solution must align with the requirements inspired by the LSEG-API-Samples/Article.EMA.CPP.StandbyServiceConsumer, supporting features like leader election, heartbeat monitoring, and failover, while integrating with the Refinitiv Real-Time server using EMA Java. Below is a list of reliable approaches, including those previously discussed (Redis, EhCache with shared file, UDP multicast, Oracle database, JGroups) and additional viable options, focusing on their mechanisms, reliability factors, and trade-offs.

Reliable Approaches for EMA Consumer Application on Two Linux Hosts

1. Redis-Based Coordination

- **Description:** Use Redis as a distributed key-value store for leader election and heartbeat management. The active host sets a key (`ema:activeHost`) using `setIfAbsent` with a TTL, and updates a heartbeat key (`ema:heartbeat:<hostId>`). The standby host monitors for stale heartbeats or missing keys to take over.
- **Implementation:**
 - Leader election: Use Redis `SETNX` (set if not exists) to designate the active host.
 - Heartbeat: Active host updates a timestamp key every 5 seconds, with a 15-second TTL.
 - Failover: Standby host takes over if the heartbeat is stale or the active host key expires.
 - EMA integration: Active host runs `OmmConsumer`, stopping on connection failure (`StatusMsg`).

- **Reliability Factors:**

- **Pros:**

- Atomic operations (`SETNX`, `TTL`) ensure single active host, preventing race conditions.
- Fast and lightweight, with low-latency key-value operations.
- Widely used, with mature libraries (e.g., Spring Data Redis).
- Redis Sentinel or Cluster provides high availability for Redis itself.

- **Cons:**

- Requires a Redis server, adding infrastructure overhead.
- Single point of failure if Redis is not configured for high availability.

- **Setup:**

- Install Redis server, configure in `application.properties` (`spring.redis.host/port`).
 - Use `StringRedisTemplate` for key-value operations.
- **Use Case:** Ideal for environments with existing Redis infrastructure, requiring fast and reliable coordination.

² EhCache with Shared File System

- **Description:** Use EhCache for local caching of session state (e.g., active/inactive status, market data) on each host, combined with a shared file system (e.g., NFS) for coordination. A lock file stores the active host's ID and heartbeat timestamp.
- **Implementation:**
 - Leader election: Hosts attempt to create/write to a lock file (e.g., /shared/ema-active-host.lock) atomically.
 - Heartbeat: Active host updates the lock file's timestamp every 5 seconds.
 - Failover: Standby host takes over if the timestamp is stale (e.g., >15 seconds) or the file is missing.
 - EMA integration: Active host runs `OmmConsumer` , caching data in EhCache, stopping on connection failure.
- **Reliability Factors:**
 - Pros:
 - No external server (e.g., Redis) required, reducing dependencies.
 - EhCache provides fast local caching for session state and market data.
 - Simple for small setups with shared storage.
 - Cons:
 - File-based coordination is prone to race conditions and file system latency.
 - Shared file system (e.g., NFS) introduces a single point of failure.
 - Less robust than distributed systems like Redis or ZooKeeper.
- **Setup:**
 - Configure a shared file system accessible to both hosts.
 - Use EhCache dependency (`org.ehcache:ehcache`) for local caching.
- **Use Case:** Suitable for environments with a reliable shared file system but no Redis or database infrastructure.

³ UDP Multicast-Based Coordination

- **Description:** Use UDP multicast for direct host-to-host communication, where the active host sends periodic heartbeats to a multicast group, and the standby host monitors for missed heartbeats to take over.
- **Implementation:**
 - Leader election: First host to send heartbeats becomes active.
 - Heartbeat: Active host sends `hostId:timestamp` to a multicast group (e.g., `239.255.0.1:5000`) every 5 seconds.
 - Failover: Standby host takes over if no heartbeats are received for 15 seconds or on connection failure.
 - EMA integration: Active host runs `OmmConsumer`, stopping on `StatusMsg` indicating connection loss.
- **Reliability Factors:**
 - Pros:
 - No external dependencies (e.g., Redis, database, or file system).
 - Lightweight network-based communication.
 - Simple for two-host setups.
 - Cons:
 - Multicast packets can be lost in unreliable networks, causing false failovers.
 - Requires network support for multicast, which may need router configuration.
 - Less robust for larger clusters compared to JGroups or ZooKeeper.
- **Setup:**
 - Configure multicast group and port in `application.properties`.
 - Use Java's `MulticastSocket` for communication.
- **Use Case:** Good for environments without external infrastructure but with reliable multicast-enabled networks.

4 Oracle Database Coordination

- **Description:** Use an Oracle database to store the active host's ID and heartbeat timestamp in a table (`EMA_ACTIVE_HOST`). Hosts use transactional updates to manage leadership.
- **Implementation:**
 - Leader election: Hosts attempt to update the table with their `hostId` and timestamp, using a query that checks the current host or stale timestamp.
 - Heartbeat: Active host updates the timestamp every 5 seconds.
 - Failover: Standby host takes over if the heartbeat is stale (>15 seconds) or on connection failure.
 - EMA integration: Active host runs `OmmConsumer`, stopping on connection failure.
- **Reliability Factors:**
 - **Pros:**
 - Transactional updates ensure atomicity, preventing race conditions.
 - Leverages existing Oracle infrastructure, with high availability via Oracle RAC.
 - Persistent state allows auditing of host changes.
 - **Cons:**
 - Database operations are slower than Redis or JGroups.
 - Requires Oracle database setup and maintenance, adding complexity.
 - Database connectivity issues could delay failover.
- **Setup:**
 - Create `EMA_ACTIVE_HOST` table in Oracle.
 - Use Spring Data JPA with `ojdbc8` driver for database access.
- **Use Case:** Ideal for environments with existing Oracle infrastructure, prioritizing data persistence and SQL familiarity.

5 JGroups-Based Coordination

- **Description:** Use JGroups for cluster coordination, leveraging its built-in leader election and failure detection. The coordinator host is active, and the other is inactive.
- **Implementation:**
 - Leader election: JGroups assigns the coordinator role based on cluster view (first member in the view).
 - Heartbeat/Failure Detection: JGroups' `FD_ALL` protocol detects host failures (~10 seconds).
 - Failover: Non-coordinator takes over when JGroups signals a new coordinator due to failure or connection loss.
 - EMA integration: Coordinator runs `OmmConsumer`, supporting `ChannelSet` for primary/standby servers, stopping on connection failure.
- **Reliability Factors:**
 - **Pros:**
 - No external dependencies (e.g., Redis, database), runs within the JVM.
 - Robust failure detection and leader election with `FD_ALL` and `GMS` protocols.
 - Supports UDP or TCP, adaptable to network constraints.
 - Scales to more than two hosts if needed.
 - **Cons:**
 - Requires network configuration (multicast or TCP).
 - Slightly more complex than raw UDP multicast due to JGroups configuration.
- **Setup:**
 - Include `jgroups` dependency and configure `udp.xml` (or `tcp.xml`).
 - Use JGroups `JChannel` for cluster management.
- **Use Case:** Best for environments needing robust, dependency-free coordination with flexible network options.

⁶ TCP-Based Heartbeat

- **Description:** Use direct TCP sockets for heartbeat communication between hosts. The active host listens on a port, and the standby pings it to detect failures.

- **Implementation:**

- Leader election: First host to start a TCP server becomes active.
- Heartbeat: Standby host pings the active host's TCP port every 5 seconds.
- Failover: Standby takes over if pings fail for 15 seconds or on connection failure.
- EMA integration: Active host runs `OmmConsumer`, stopping on `StatusMsg` indicating connection loss.

- **Reliability Factors:**

- **Pros:**
 - Simple implementation using Java's `ServerSocket` and `Socket`.
 - No external dependencies or complex network setup.
- **Cons:**
 - TCP connections can be brittle in unreliable networks.
 - Manual heartbeat logic is less robust than JGroups or Redis.
 - Firewall configuration required for TCP ports.
- **Setup:**
 - Configure TCP port in `application.properties`.
 - Implement `ServerSocket` for the active host and `Socket` for the standby.
- **Use Case:** Suitable for simple setups with stable networks and minimal infrastructure.

7 ZooKeeper-Based Coordination

- **Description:** Use Apache ZooKeeper for distributed coordination, leveraging its ephemeral nodes for leader election and heartbeat monitoring.
- **Implementation:**
 - Leader election: Hosts compete to create an ephemeral znode (e.g., `/ema/activeHost`). The successful host becomes active.
 - Heartbeat: Active host updates a timestamp znode every 5 seconds.
 - Failover: Standby watches the znode; takes over if it disappears or the heartbeat is stale (>15 seconds).
 - EMA integration: Active host runs `OmmConsumerX`, stopping on connection failure.
- **Reliability Factors:**
 - **Pros:**
 - Highly reliable with ZooKeeper's consensus-based coordination.
 - Ephemeral znodes automatically handle host failures.
 - Scales well for larger clusters.
 - **Cons:**
 - Requires ZooKeeper server setup, adding infrastructure overhead.
 - More complex than Redis or JGroups for small setups.
- **Setup:**
 - Install ZooKeeper server, configure in `application.properties`.
 - Use Apache Curator for simplified ZooKeeper integration.
- **Use Case:** Ideal for environments needing robust, scalable coordination with existing ZooKeeper infrastructure.

8 Spring Integration with Shared Message Channel

- **Description:** Use Spring Integration to create a shared message channel (e.g., via JMS or AMQP) for coordination. Hosts publish and subscribe to messages for leader election and heartbeats.
- **Implementation:**
 - Leader election: Hosts publish their intent to become active; a locking mechanism (e.g., via message queue) ensures one winner.
 - Heartbeat: Active host sends periodic messages to the channel.
 - Failover: Standby takes over if heartbeats stop or on connection failure.
 - EMA integration: Active host runs `OmmConsumer`, stopping on `StatusMsg`.
- **Reliability Factors:**
 - **Pros:**
 - Integrates well with Spring Boot ecosystem.
 - Leverages existing message brokers (e.g., RabbitMQ, ActiveMQ).
 - Flexible for complex messaging patterns.
 - **Cons:**
 - Requires a message broker, adding infrastructure complexity.
 - Higher latency than Redis or JGroups due to message queue overhead.
- **Setup:**
 - Configure a message broker (e.g., RabbitMQ) and Spring Integration dependencies.
 - Define a shared channel for coordination messages.
- **Use Case:** Suitable for environments with existing message broker infrastructure and Spring expertise.

Common Features Across Approaches

- **EMA Integration:** All approaches use `OmmConsumerClient` to handle `RefreshMsg`, `UpdateMsg`, and `StatusMsg`, with failover triggered by non-open `StatusMsg`.
- **Primary/Standby Servers:** Can be supported using EMA's `ChannelSet` (e.g., `Channel_Active`, `Channel_Standby`), as shown in the JGroups implementation.
- **Failover Triggers:** Connection loss (`StatusMsg`) or host failure (detected via heartbeat or cluster view).
- **Spring Boot:** Provides configuration (`application.properties`), dependency injection, and lifecycle management (`@PostConstruct`, `@PreDestroy`).

Reliability Considerations

- **Atomicity:** Redis, Oracle, and ZooKeeper provide strong consistency for leader election. JGroups and Spring Integration are also reliable but depend on network stability. File-based and TCP/multicast approaches risk race conditions.
- **Failure Detection:** JGroups (`FD_ALL`), ZooKeeper (ephemeral znodes), and Redis (TTL) offer robust failure detection. TCP/multicast and file-based approaches are less reliable due to network or I/O issues.
- **Infrastructure Overhead:** Redis, Oracle, ZooKeeper, and Spring Integration require external servers. JGroups, TCP, and multicast are dependency-free within the JVM. File-based requires a shared file system.
- **Scalability:** JGroups and ZooKeeper scale well to more hosts. Redis and Oracle are viable but less flexible. TCP, multicast, and file-based approaches are best for two hosts.
- **Network Dependency:** Multicast, JGroups (UDP), and TCP require network configuration. Redis, Oracle, and ZooKeeper need stable server connectivity.
- **Performance:** Redis and JGroups are fastest due to low-latency operations. Oracle and Spring Integration are slower due to database or message queue overhead.

Recommended Approach

- **JGroups**: Recommended for its balance of reliability, simplicity, and no external dependencies. It provides robust leader election and failure detection, supports UDP or TCP, and integrates well with Spring Boot. The `ChannelSet` support for primary/standby servers ensures comprehensive failover.
- **Redis**: A close second if you have Redis infrastructure and prioritize low-latency coordination.
- **ZooKeeper**: Best for larger clusters or environments with existing ZooKeeper setups.
- **Oracle**: Suitable if you already use Oracle and need persistent state.
- **Others**: TCP/multicast for minimal setups with stable networks; Spring Integration for message-driven architectures; file-based only if no other options are feasible.

Adding Primary/Standby Servers

For any approach, add primary/standby server support by configuring a `ChannelSet` in `OmmConsumerConfig`:

```
java
config.addChannel("Channel_Active", Map.create()
    .add(MapEntry.create("ChannelType", OmmConsumerConfig.ChannelTypeEnum
        .add(MapEntry.create("Host", activeHost))
        .add(MapEntry.create("Port", activePort))));

config.addChannel("Channel_Standby", Map.create()
    .add(MapEntry.create("ChannelType", OmmConsumerConfig.ChannelTypeEnum
        .add(MapEntry.create("Host", standbyHost))
        .add(MapEntry.create("Port", standbyPort))));

config.addChannelSet("ChannelSet_1", new String[]{"Channel_Active", "Channel_Standby"});
```

Use `channelSet("ChannelSet_1")` in `startConsumer`. Update `application.properties` with `ema.active.host/port` and `ema.standby.host/port`.

Production Considerations

- **Authentication:** Add Refinitiv credentials to `OmmConsumerConfig` if required.
- **Logging:** Use SLF4J for detailed logging in production.
- **Monitoring:** Monitor cluster health (JGroups, ZooKeeper) or server availability (Redis, Oracle).
- **Tuning:** Adjust heartbeat intervals (e.g., 5s/15s) and dispatch timeouts (e.g., 1s) for performance.
- **Testing:** Simulate host and server failures to verify failover.

Each approach can be implemented using the Spring Boot EMA application framework, with JGroups being the most versatile for most environments. Refer to the LSEG Developer Community for EMA Java details and respective documentation for each technology (e.g., JGroups, Redis, ZooKeeper).

If you need a detailed implementation for any of these approaches or specific tweaks (e.g., TCP-based JGroups, authentication), please let me know!